

CSE 422: Artificial Intelligence

Logistic Regression

Swakkhar Shatabda

BRAC University

December 2, 2024

Contents

- 1 Classification
- 2 Sigmoid
- 3 Cross Entropy
- 4 Stochastic Gradient Descent

Classification

- 1 In classification problems, we are given data as X and labels as y .
- 2 Here, we are upto learn a model where, y will be predicted as a function of X .
- 3 In classification, the label y is categorical or discrete in value.
- 4 For example, suppose you are given many features of a fish, like length, weight, eggs, months and you have to predict whether it is legal to be caught or not. This problem can be formulated as a classification problem.

Data

Here is how data looks like in a supervised setting:

| instance no | features | | | | label |
|-------------|----------|----------|----------|----------|----------|
| | x_1 | x_2 | x_3 | x_4 | y |
| | length | weight | has eggs | month | legal? |
| 1 | 10 | 250 | 1 | 12 | No |
| 2 | 20 | 1250 | 0 | 1 | Yes |
| 3 | 15 | 750 | 1 | 2 | No |
| \vdots | \vdots | \vdots | \vdots | \vdots | \vdots |
| m | 17 | 550 | 0 | 3 | Yes |


Experiments

We will first try to predict the class of the dataset based on two features, x_1 and x_2 .

| | |
|----|-------------------------|
| 1 | -0.017612, 14.053064, 0 |
| 2 | -1.395634, 4.662541, 1 |
| 3 | -0.752157, 6.53862, 0 |
| 4 | -1.322371, 7.152853, 0 |
| 5 | 0.423363, 11.054677, 0 |
| 6 | 0.406704, 7.067335, 1 |
| 7 | 0.667394, 12.741452, 0 |
| 8 | -2.46015, 6.866805, 1 |
| 9 | 0.569411, 9.548755, 0 |
| 10 | -0.026632, 10.427743, 0 |
| 11 | 0.850433, 6.920334, 1 |
| 12 | 1.347183, 13.1755, 0 |
| 13 | 1.176813, 3.16702, 1 |
| 14 | -1.781871, 9.097953, 0 |
| 15 | -0.566606, 5.749003, 1 |
| 16 | 0.931635, 1.589505, 1 |
| 17 | -0.024205, 6.151823, 1 |
| 18 | -0.036453, 2.690988, 1 |
| 19 | -0.196949, 0.444165, 1 |
| 20 | 1.014459, 5.754399, 1 |
| 21 | 1.985298, 3.230619, 1 |
| 22 | -1.693453, -0.55754, 1 |

Experiments

We will first try to predict the class of the dataset based on two features, x_1 and x_2 .



Logistic Regression

At first, we are going to try a linear classifier called logistic regression. We can apply logistic regression when the data is linearly separable.

- The relationship will be predicted as:

$$y = w_0 + w_1 x_1$$

- This is again an equation of a straight line
- We need the best line that separates blue from the orange
- learn w_0, w_1, \dots
- Can we use gradient descent here? A little trick required!


Gradient Descent for Logistic Regression

- The cost function / loss function of gradient descent

$$e = \frac{1}{2} \sum_{i=1}^m (\hat{y}(i) - y(i))^2$$

- This time too predicted label \hat{y} is a function of \vec{x} and w
- The labels are discrete, for this binary classification two labels 0 (no or negative) and 1 (yes or positive)
- Now, we try to define \hat{y} with help of the weights or coefficients of the line.

Linear Classification



- This linear classifier divides instances based on the local wrt the line, on the right positive, negative on the left
- Any point on the line satisfies the equation. Any point on the right (3,1) yields positive result and any point on the left (1,1) yields negative result.
- Based on this we can define a linear classifier

Linear Classification

This following function will help us in making decision:


$$f(\vec{x}) = w_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n$$

LinearClassifier

```
1  if  $f(\vec{x}) > 0$  or  $w_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n > 0$ 
2      return 1
3  else return 0
```


This simple classifier just checks whether a point is on the left or right.

A step function!



Alas! This is not a continuous function and thus not differentiable. We can't calculate gradients! We need to find an alternate!

A sigmoid function!



$$\sigma(\vec{x}) = \frac{1}{1 + \exp(-\vec{x})}$$

Good things about sigmoid!

- 1 Its continuous and differentiable.
- 2 $\sigma'(\vec{x}) = \sigma(\vec{x})(1 - \sigma(\vec{x}))$

Lets go back to the loss function now.

A new loss function - Cross-entropy


Cross-entropy loss, or log loss, measures the performance of a classification model whose output is a probability value between 0 and 1.

$$e = \sum_{i=1}^m (-y(i)\log(\hat{y}(i)) - (1 - y)\log(1 - \hat{y}))$$


Cross Entropy Loss Function

How it works?

Loss function for true label $y=1$



Loss function for true label $y=0$



$$e = \sum_{i=1}^m (-y(i) \log(\hat{y}(i)) - (1 - y) \log(1 - \hat{y}))$$

Cross Entropy Loss Function

How to find the gradient? Lets try!

$$\begin{aligned}\frac{\delta e}{\delta w_0} &= \frac{\delta}{\delta w_0} \sum_{i=1}^m (-y(i) \log(\hat{y}(i)) - (1 - y) \log(1 - \hat{y}(i))) \\ &= \sum_{i=1}^m (-y(i) \frac{1}{\hat{y}(i)} \hat{y}(i)(1 - \hat{y}(i)).1 - (1 - y) \frac{1}{(1 - \hat{y}(i))} (-1) \hat{y}(i)(1 - \hat{y}(i)).1) \\ &= \sum_{i=1}^m (-y(i) + y(i) \hat{y}(i) + \hat{y}(i) - y(i) \hat{y}(i)).1 \\ &= \sum_{i=1}^m (\hat{y}(i) - y(i)).1\end{aligned}\tag{1}$$

In a similar way,

$$\frac{\delta e}{\delta w_i} = \sum_{i=1}^m (\hat{y}(i) - y(i)).x_i\tag{2}$$

Now the same gradient descent will work!

Comments on Gradient Descent

- 1 Slow when the dataset is too large!
- 2 Rather learning the whole dataset, possible to learn in chunks!
- 3 What if we process only 1 single item at each iteration?
- 4 Lets have another look!

Gradient Descent Algorithm

GRADIENTDESCENT($X, y, \alpha, \text{maxIter}$)

```
1  for  $j = 1$  to  $m$ 
2       $x_0(j) = 1$ 
3   $w_0, w_1, \dots, w_n$  initialized randomly
4   $iter = 0$ 
5  while  $iter++ \leq \text{maxIter}$ 
6      for  $j = 0$  to  $n$ 
7           $slope_j = 0$ 
8      for  $i = 1$  to  $m$ 
9           $\hat{y} = w_0 + w_1x_1(i) + w_2x_2(i) + \dots + w_nx_n(i)$ 
10          $e = \hat{y} - y(i)$ 
11         for  $j = 0$  to  $n$ 
12              $slope_j = slope_j + e \times x_j(i)$ 
13     for  $j = 0$  to  $n$ 
14          $w_j = w_j - \alpha \times slope_j$ 
15 return  $w_0, w_1, \dots, w_n$ 
```

Lighter Gradient Descent Algorithm

LIGHTERGRADIENTDESCENT($X, y, \alpha, \text{maxIter}$)

```
1  for  $j = 1$  to  $m$ 
2       $x_0(j) = 1$ 
3   $w_0, w_1, \dots, w_n$  initialized randomly
4   $iter = 0$ 
5  while  $iter++ \leq \text{maxIter}$ 
6      for  $j = 0$  to  $n$ 
7           $slope_j = 0$ 
8           $i = iter$ 
9           $\hat{y} = w_0 + w_1x_1(i) + w_2x_2(i) + \dots + w_nx_n(i)$ 
10          $e = \hat{y} - y(i)$ 
11         for  $j = 0$  to  $n$ 
12              $slope_j = slope_j + e \times x_j(i)$ 
13         for  $j = 0$  to  $n$ 
14              $w_j = w_j - \alpha \times slope_j$ 
15 return  $w_0, w_1, \dots, w_n$ 
```

Thats it!

Thank you